

REZA RAJAN

Software Engineer - Data Systems

📞 1-407-649-0447
✉ reزارajan@gmail.com
🌐 /in/reزارajan
🔗 reزارajan

Summary

Software engineer specializing in data systems, having started in full-stack development and progressed into building end-to-end data platforms, from lakehouse architectures to their infrastructure-as-code foundations.

Core Competencies

Data Engineering: ETL/ELT, CDC (CDF, Debezium, Airbyte), Data Lakehouse & Modeling, Databricks, Lakebase, Trino, Apache Airflow, dbt, Spark, PostgreSQL, AWS (Glue, Redshift, Step Functions, Lambda)

Software Engineering: Python, Golang, SQL, Bash, Git, CI/CD (GitHub Actions)

Infrastructure as Code: Kubernetes, Docker, Terraform/OpenTofu, Ansible, Nix, AWS, Azure

Analytics & Visualization: Power BI, Grafana, Looker Studio

Experience

Ministry of Education Trinidad and Tobago - Data Specialist

2026 - Present

- Founded the Ministry's first data capability, winning executive approval to automate reporting across all 31 ministry services by delivering a serverless Azure Databricks lakehouse proof of concept.
- Established data governance with full lineage and provenance coverage of ingested sources via Unity Catalog.

Mochi - Co-Founder / Infrastructure Engineer

Dec 2022 - Present

- Sustained 98% Cardano block-production uptime with a Python Kubernetes sidecar for replica failover.
- Cut infrastructure costs ~70% vs. cloud by deploying on-prem after a cost/reliability trade-off analysis.
- Cut rollouts from 3+ hours to <15 minutes with declarative Talos Linux IaC and GitHub Actions CI/CD.
- Reduced incident response time by ~60% with Prometheus and Grafana monitoring and alerting.

myKaarma - Data Engineer

Feb 2022 - Jan 2024

- Delivered automated reporting to 5,000+ dealerships for 3 enterprise clients through AWS Step Functions orchestration.
- Cut daily report generation query times by up to 50% through MariaDB stored procedure indexing.
- Initiated the lakehouse modernization after hitting MariaDB scalability limits, winning buildout approval with a PySpark/Redshift proof of concept that cut terabyte-scale data processing time by ~70%.
- Stood up the approved buildout with the Principal Engineer, delivering its architecture, roadmap, and first ingestion pipelines.

Education

Accreditation in CS (DSA), UC Berkeley

2025

Bsc. Mechanical Engineering, University of Waterloo

2015 - 2020

Certifications

AWS Certified Data Engineer, Amazon

In-Progress

Project Initiation, Google

2026

Foundations of Project Management, Google

2026

Data Engineer, Udacity

2024

Self-Driving Car Engineer, Udacity

2021

Data Analyst, Udacity

2020